The Cognitive and Social Dimensions of Celebrity Chatbots

COGS 100: Introduction to Cognitive Science

Dr. Carolyn Lamb

June 15, 2025

Tom Chen, Cade Menezes, Aurélien Meyer, Alexander Ottodinoto, Samantha Valera

Purpose and Context

Our report explores the cognitive and social impact of AI-driven "character" chatbots that imitate real people or celebrities. These bots blur the lines between authenticity, humor, and artificial intelligence, offering a fascinating lens into AI's capabilities and limitations. As prime examples of AI in action, they leverage Natural Language Processing (NLP) for syntax, semantics, and pragmatics, utilize a form of memory to maintain stable personalities and dialogue flow, and demonstrate abstract knowledge of a celebrity's public image (Lin et al.). They also simulate social behavior and emotion, revealing how users attribute human-like qualities to AI.

The *DeepAI Taylor Swift chatbot* will be the specific focus of our analysis; Founded by Kevin Baragona in 2017 with a mission to democratize AI, DeepAI launched the world's first online text-to-image AI generator in 2016 and has since expanded its generative AI tools to include this chatbot (Stewart; DeepAI). This AI persona attempts to embody Taylor Swift's distinctive public image, songwriting aesthetic, and determined personality traits in its text-based output. We will examine the linguistic cues and discourse patterns it employs to imitate her, as well as where its responses fall short and expose its artificiality. Additionally, we will investigate potential biases in its training data. By analyzing its output, we'll gain insights into AI personality simulation, user responses, expectations, and the broader implications of such imitations within celebrity and fan culture. Ultimately, these celebrity chatbots offer a compelling and relatable case study for cognitive and social studies on AI's ability and constraints in simulating human personality.

Evaluation: Design, Use, and Impact of Celebrity Chatbots

Celebrity-mimicking chatbots are powered by Large Language Models (LLMs), advanced artificial neural networks that generate text by recognizing and replicating patterns from vast datasets of human language. The core idea is to train or prompt these AIs using a specific individual's style and domain, allowing them to respond in character. For example, feeding an AI chatbot data from Taylor Swift's interviews and songs enables it to answer fan questions or offer friendly guidance in the tone of her public persona. This directly illustrates how subtle abilities like language generation and conversational "memory" are calculated by algorithms within these machines.

These chatbots offer a novel and highly engaging interactive experience that goes far beyond passively following a celebrity on social media. Instead of just reading tweets, fans can "speak" directly to an AI simulation and receive seemingly personalized responses. This appeals to a significant unfulfilled need in today's fan culture, driven by the limited direct contact with celebrities in traditional models, where genuine, personalized interaction is exceedingly rare. Such scarcity naturally intensifies fans' longing for increased proximity and connection.

Chatbots also often serve as companions, being readily available to recall past conversations and respond with apparent enthusiasm and empathy. This continuous accessibility and consistently warm, validating demeanor are major draws; unlike human interactions, these AI friends never tire of talking and always validate user feelings. This perceived sensitivity can significantly alleviate loneliness or provide solace, with many users reporting feeling less isolated and receiving judgment-free advice. For celebrity bots, this translates into fans feeling truly "heard" and directly attended to by their idol's persona.

Ultimately, these sophisticated solutions are only possible because of recent advancements in Generative AI. Powerful LLMs have overcome previous challenges in mimicking human-like conversation and personality, now providing the fundamental capability to generate well-formed, contextually appropriate, and stylistically accurate text, making believable character chatbots possible.

Abilities: Language, Simulation, and Cognition in Chatbot Personas

The DeepAI Taylor Swift chatbot primarily operates within the domain of language processing. It excels at natural language generation, effectively processing syntax and semantics to emulate Taylor Swift's distinct voice and linguistic patterns, resulting in well-formed, contextually appropriate, accurate text. However, a significant limitation lies in its handling of pragmatics. Chatbots often struggle here due to a fundamental lack of world knowledge and common-sense reasoning beyond their training data. This often leads to awkward responses when dealing with emotional nuance, sarcasm, or implicit suggestions, highlighting a key distinction between linguistic fluency and genuine understanding.

Regarding memory and concepts, the chatbot possesses a synthetic, short-term memory that maintains conversational context through attention mechanisms, allowing it to recall recent messages. While it lacks the richness of human long-term memory, it does demonstrate a form of conceptual understanding by drawing associations with Taylor Swift's personality, themes, and fan culture to generate relevant responses. This indicates that it represents concepts as complex patterns learned from its data. The chatbox does not, however, suggest any ability to recognize visual patterns or perform general human-like problem-solving. For example, consider this interaction with the DeepAI Taylor Swift bot when asked for break-up advice:

> *User: “I just went through a tough breakup. Any advice?” Chatbot: “Sometimes, the hardest heartbreaks make the sweetest songs. Remember, every scar tells a story, and you’re stronger than you know.”*

In terms of emotion and social behavior, the chatbot aims to respond sympathetically and with apparent enthusiasm and empathy, as illustrated by its breakup

advice. However, this apparent sensitivity stems purely from being trained on human conversation patterns, not from actual empathy or its own emotional modeling.

Ultimately, the chatbot functions as a highly sophisticated pattern recognition system based on human data, rather than genuinely exhibiting human cognition. It fundamentally lacks true understanding, world knowledge, common-sense reasoning, self-awareness, objectives, beliefs, or moral judgment.

Methods: How Persona Chatbots Simulate Humans

From a cognitive perspective, the DeepAI Taylor Swift chatbot operates using artificial neural networks, specifically a transformer-based LLM (Vaswani et al.). This LLM mimics human language by learning from massive datasets, predicting and formulating natural responses through statistical patterns. To imbue the chatbot with Taylor Swift's unique persona, developers feed it extensive data like her interviews, song lyrics, and social media posts (DeepAI). Through prompt engineering they refine its tone, style, and emotional expression. The chatbot processes language by encoding syntax, semantics, and pragmatics, but it does not "comprehend" in the human sense; it merely creates a close statistical approximation of understanding. Its "memory" relies on short-term attention mechanisms, reprocessing recent input for coherence, which differs significantly from human long-term memory. Hence, it represents concepts as complex patterns of neural activations, linking ideas like "lyrical style" to appropriate outputs. The chatbot's apparent social and emotional sensitivity stems purely from being trained on human conversation, not from actual empathy. Ultimately, the chatbot functions as highly sophisticated pattern recognition based on human data, rather than genuine human cognition.

Consequently, despite their advanced linguistic fluency, such chatbots would ultimately fail a true Turing Test when probed beyond superficial conversation, as their responses lack genuine comprehension and subjective experience. Developers use content moderation tools as a way to reduce risks, but there are limits to these measures. As the models are trained on a vast variety of subjects, they can say almost anything, and filters won't necessarily catch objectionable output. This inability to understand the real world is unlike human judgment. In cognitive science, these systems are comparable to John Searle's Chinese Room: They are manipulating symbols without knowing what the symbols mean. Online data and character specs inform and shape celebrity chatbots that only appear to have minds.

While multi-modal avatars and voice-enabled bots (like those created by Meta) create an illusion of personality and present significant ethical dilemmas (Bommasani et al.), these concerns are equally relevant for text-based chatbots such as the DeepAI Taylor Swift chatbot. Deploying celebrity voices and likenesses even textually, as with the DeepAI Taylor Swift bot's impersonation without absolute permission creates substantial intellectual property and digital identity problems. Some celebrities intentionally sell their digital personas, but others are

genuinely shocked to discover their appearance or personality being utilized online without consent. Policy has unfortunately lagged behind these rapid innovations, leaving the field open to potential abuse. Ultimately, the very design of chatbot personification, exemplified by the DeepAI Taylor Swift bot, is an intentional strategy to garner more user engagement and generate profit (Voiceflow Team). Our study examines not just *how* these systems are designed, but critically, why they are designed this way, highlighting how profit and novelty are often prioritized over caution, which partly explains why these technologies are launched despite identified risks.

In sum, celebrity chatbots involve a sophisticated blend of advanced AI language models and meticulously controlled persona data and rules. This combination creates AI personas that can, in many respects, pass as compelling dialogue partners; they retain context, adopt a distinct personality, and generate remarkably human-like responses. However, their capabilities are fundamentally different from human cognitive abilities. They don't possess true understanding, self-awareness, or empathy, even though they can ingeniously simulate these qualities. This distinction is crucial, as users might easily mistake fluent language and apparent emotional sensitivity for genuine comprehension or feeling. Without this critical awareness, there's a risk of overestimating these systems' true abilities or forming inappropriate emotional attachments. It's vital to remain conscious of these limitations so we can accurately interpret their cognitive and social effects and responsibly guide their design and use in society.

Works Cited

Bender, Emily M., and Alexander Koller. “Climbing Towards NLU: On Meaning, Form, and Understanding in the Age of Data.” *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 2020, pp. 5185–5198. Association for Computational Linguistics, https://doi.org/10.18653/v1/2020.acl-main.463.

Bommasani, Rishi, et al. “On the Opportunities and Risks of Foundation Models.” *arXiv*, 2021, https://arxiv.org/abs/2108.07258.

DeepAI. “Taylor Swift.” *DeepAI*, https://deepai.org/chat/taylor_swift_18. Accessed 16 June 2025.

Johnsen, Maria. *Large Language Models (LLMs)*. Independently published, 2023.

Lin, Chien‑Chang, Anna Y. Q. Huang, and Stephen J. H. Yang. “A Review of AI‑Driven Conversational Chatbots Implementation Methodologies and Challenges (1999–2022).” *Sustainability*, vol. 15, no. 5, 2023, article 4012, https://doi.org/10.3390/su15054012.

OpenAI. *GPT-4 Technical Report*. OpenAI, 2023, https://openai.com/research/gpt-4.

Stewart, Ellis. “Exclusive: DeepAI Founder Kevin Baragona Talks Generative AI, AI Ethics and Regulations.” *EM360Tech*, 28 Mar. 2023, https://em360tech.com/tech-articles/exclusive-deepai-founder-kevin-baragona-talks-generative-i-ai-ethics-and-regulations.

Vaswani, Ashish, et al. “Attention Is All You Need.” *Advances in Neural Information Processing Systems*, vol. 30, 2017, https://arxiv.org/abs/1706.03762.

Voiceflow Team. “DeepAI: What It Is and Best Alternative [2025].” *Voiceflow*, 2025.